

Sharon Dev Saseendran

Buffalo, NY | [portfolio](#) | sharon.devs99@gmail.com | (716) 573 9185
linkedin.com/in/sharon-dev-s-8897a659/ | github.com/sharondevs

SKILLS

Languages: TypeScript · Go · Java · Dart · Python · SQL **MLOps:** PyTorch · vLLM · LangChain · Pinecone · LlamaIndex
Frameworks: Next.js · Golang · FastAPI · Flutter · N8N · Flask · Spring Boot · PostgreSQL · MSSQL · Elasticsearch · Apache Kafka · Redis
Observability: Azure (Container Apps, Functions, Service Bus, App Insights, Elastic Pool) · GCP · GitHub Actions · Prometheus · Grafana · JUnit

WORK EXPERIENCE

OXmaint Inc

Lead Software Engineer

July 2025 – Present | Sunnyvale, CA

Next.js | TypeScript | Golang | n8n | Flutter | Kafka | Redis

- Owned service architecture for OXmaint's AI-enabled EAM/CMMS platform with a **Next.js/TypeScript** frontend and **Golang** microservices backend, SQL schema migrations, Redis caching and multi-tenant auth; migrated 12+ legacy workflows, increasing daily active usage by 40%.
- Closed 100% of audit-identified XSS vectors** by hardening the frontend with nonce-based CSP and CSRF rotation; contributed to GitHub Actions CI/CD pipeline with Docker staging, Prometheus + Grafana alerting (P95 latency, error rates), and structured logging.
- Improved multilingual onboarding** by shipping i18n-based locale routing with middleware-driven auto-locale detection and per-locale SSG.
- Reduced customer onboarding effort by 40%** by shipping 3 OxyAI agentic workflows on n8n — org-scoped context retrieval, tool-calling over internal APIs, LLM routing (GPT-OSS 120B) with automatic fallback, and Redis-backed session state with token-budget windowing.
- Cut sensor setup time by 60%** by engineering the end-to-end IoT ingestion pipeline — ChirpStack sensor networks, MQTT services (Eclipse Paho Go), automatic device discovery, tenant-scoped onboarding, and downstream analytics triggers for predictive maintenance.
- Worked on third-party enterprise integrations with QuickBooks (REST/OAuth 2.0) and SAP BTP (OData V2) — field-level mapping for work orders, purchase orders, and billable costs with automated invoice generation and PO reconciliation.
- Drove the end-to-end customer journey** — from prospect demos and technical discovery through solution architecture, onboarding, and go-live; managed CSRs to ensure **retention and expansion**.
- Owned the distribution build of the OXmaint AI Flutter app**, shipping production CMMS features with modular MVVM architecture, secure OAuth2, Firebase notifications; built on-device TTS, **wake-word detection**, and **Meta wearable (MWDAT SDK)** integration to enable hands-free task execution, **improving field-user turnaround by 30%**, expanding voice-assisted coverage across major CMMS workflows.

Moonraft Innovation Labs (acquired by UST)

February 2023 - July 2024 | Bangalore, India

Developer 1

Spring Boot | Flask | Flutter | FlutterFlow | Elasticsearch | Pinecone | ETL | MongoDB

- Delivered a cross-platform **Flutter** app for a UK mortgage client, scaling MVP to production for **65k+** users and achieving **85%** crash-free sessions via **GA** and **FlutterFire** profiling; build splash-screen transitions with **AnimatedBuilder**, **SVG masks**, **ClipPath**.
- Plug-and-play semantic-search module:** Spring Boot/FastAPI + Crawler4J(web crawling), Apache Kafka job queue monitoring, HuggingFace embeddings indexed in Elasticsearch k-NN. Delivered a React portal (WebSockets/long-polling, audit logs, alerts) orchestrating **20+ concurrent scrape/index runs**; templated configs enabled partner onboarding in **2–3 days**.
- Worked on a Product Comparison Engine catering ecommerce platforms, leading a 3-member team to orchestrate ETL pipeline to **scrape 50k+** product listings across multiple brands at scale, generating text & image embeddings to perform cosine similarity search.

Microsoft

August 2021 - February 2023 | Hyderabad, India

Cloud Solution Architect

Microsoft Intune | Azure AD | MECM | Microsoft Defender

- Managed 70+ enterprise accounts** as SME for Microsoft Intune, Azure AD, and MECM — led customer onboarding, change management, and endpoint security strategy across **200K+ managed endpoints**.
- Trained 100+ engineers** through 3–5 day workshops on Azure, Azure AD, Intune, and Defender 365; delivered technical advisory and built solution POCs for premier customers on M365 and Azure, **boosting client-side resolution rate by 15%**.

EDUCATION

University at Buffalo, State University of New York (CGPA: 3.74 / 4)

New York, United States

Masters in Computer Science (Specialization in AI/ML)

August 2024 - December 2025

Relevant Coursework: Machine Learning · Deep Learning · Data Modeling and Query Language · Computer Vision & Image Processing

Cochin University of Science and Technology, School of Engineering (CGPA: 9.36 / 10)

Kerala, India

Bachelor of Technology in Electronics and Communication

June 2017 - May 2021

PERSONAL PROJECTS

ECHO-SRE (Agentic SRE Copilot)

2026

- Built an agentic SRE copilot that root-causes production incidents via a ReAct agent over the Model Context Protocol (JSON-RPC/stdio), querying alerts, metrics, service topology, logs, and BM25-retrieved runbooks through a FastMCP server reusable by any MCP client (Claude Desktop/Cursor).
- Engineered a provider-agnostic inference gateway on the OpenAI Chat Completions API unifying local vLLM and hosted APIs (Gemini/OpenAI), with a priority-ordered chain, automatic multi-model failover (retryable-vs-fatal classification), exponential backoff, and per-provider token budgeting — removing vendor lock-in and chaining quotas (~560 req/day).
- Shipped a FastAPI Server-Sent-Events streaming service driving a React/TypeScript “SRE mode” with three per-request backends (bundled synthetic, user-uploaded scenario, or live Prometheus/Loki); instrumented with Prometheus, containerized with Docker, and CI/CD-deployed to Azure Container Apps via test-gated GitHub Actions, with a key-free offline demo and hermetic tests.

ECHO (Engaged Conversation with Human Oversight) Chat

2025

- Developed a retrieval augmented generation backend allowing Q&A based a user's profile or any documents, implementing real-time streaming responses with session-based conversation memory using Google Gemini 2.0 Flash and llamaindex.
- Built prod-ready FastAPI backend with LlamaIndex RAG framework, Google GenAI text-embedding-004 for semantic search, async document processing pipeline supporting PDF/DOCX/TXT formats and deployed with Azure Container Apps.

Visual Question Answering for Real-world accessibility

2025

- Developed a VQA system for real-world images, addressing challenges like blur, and poor lighting from the VizWiz dataset.
- We combined the pre-trained CLIP ViT-B/32 vision-language encoder and a trainable lightweight multi-head classifier to generate textual answer predictions with a test acc. of approx. **70%**.